

Hyperpartisan News Detection

Bachelor Thesis

presented by
Larissa Strauch
Matriculation Number 1518629

submitted to the
Data and Web Science Group
Prof. Dr. Ponzetto
University of Mannheim

Juli 2019

Contents

1	Introduction	1
1.1	Problem Statement	1
1.2	Contribution	1
1.3	Related Work	1
2	Fundamentals	2
2.1	Term Frequency-Inverse Document Frequency	2
2.2	Word Embeddings	3
2.2.1	Word2Vec	3
2.3	SVM	4
2.4	Multinomial Naive Bayes Classifier	4
2.5	Decision Trees and Random Forest Classifier	4
2.6	Logistic Regression Classifier	5
3	Data	6
3.1	Data Description	6
3.1.1	Dataset labelled by Publisher	7
3.1.2	Dataset labelled by Article	7
3.2	Data Analysis	8
3.3	Data Preparation	9
3.3.1	Read the XML files	10
3.3.2	Filter out the important information	10
3.3.3	Merge the Groundtruth- and Training datasets into a csv file	10
3.3.4	Remove special characters and stop words	11
3.3.5	Tokenization and Stemming of the datasets	11
4	Methodology	13
4.1	Important Algorithms	13
4.2	Vector Representation of the Text	14
4.2.1	Term Frequency-Inverse Document Frequency	14
4.2.2	Word Embeddings	17
4.3	Feature Selection and/or Feature Transformation	19
4.4	Learning Algorithms	19

CONTENTS

ii

4.4.1	SVM	19
4.4.2	Multinomial Naive Bayes Classifier	19
4.4.3	Random Forest Classifier	20
4.4.4	Logistic Regression Classifier	20

List of Algorithms

1	Random Forest	5
2	Saving a trained model	13
3	Loading a trained model	14
4	pickle.dump	16

List of Code

List of Figures

3.1	Distribution of as Hyperpartisan labelled articles by publisher . . .	7
-----	---	---

List of Tables

3.1	Publsiherers who have published more than 7 Hyperpartisan Articles	9
3.2	Publishers who have published more than 10.000 Hyperpartisan articles	9
4.1	Top 10 terms by average TF-IDF weight (By Article)	17
4.2	Top 10 terms by average TF-IDF weight (By Publisher)	17
4.3	Most similar word to 'trump'	18

Glossary

CBoW Continous Bag of Words. 3, 4, 16

IDF Inverse Term Frequency. 2, 13, 14

TF Term Frequency. 2, 13, 14

TF-IDF Term Frequency-Inverse Term Frequency. vi, 2, 13, 14, 16, 17

Chapter 1

Introduction

1.1 Problem Statement

1.2 Contribution

1.3 Related Work

Chapter 2

Fundamentals

2.1 Term Frequency-Inverse Document Frequency

Term frequency (TF) is a measure that denotes how frequently a term t appears in the document d . One way to compute TF is

$$tf(t_i, d_j) = \frac{1 + \log_{10}(f_{t_i, d_j})}{1 + \log_{10}(\max_{f_{t', d_j}: t' \in d_j})}$$

where $1 + \log_{10}(f_{t_i, d_j})$ reflects how many times the term t_i appears in document d_j and $1 + \log_{10}(\max_{f_{t', d_j}: t' \in d_j})$ is the highest occurrence of any term in document d_j .

Inverse Document Frequency (IDF) points to the assumption that the informativeness of the term t is inversely proportional to the number of documents in the collection in which the term appears.

$$idf(t_i) = \log_{10}\left(\frac{|D|}{|d' \in D : t_i \in d'|}\right)$$

Where $|D|$ is the total amount of documents in a document set and $|d' \in D : t_i \in d'|$ is the amount how many times the term t_i appears in the document set.

To compute the weight for the term t_i within the document d_j we simply multiply the *TF* and *IDF* components:

$$w_{ij} = tf(t_i, d_j) \cdot idf(t_i)$$

Therefore, TF-IDF indicates how important a word is to a document in a collection or corpus. It is often used as a weighting factor in Information Retrieval and Text Mining. The TF-IDF value increases proportionally to the number of times a word appears in the document, but is offset by the frequency of the word in the corpus, which helps to control the fact that some words are usually more common than others. (umschreiben?)

TF-IDF is easy to compute. In addition, it is possible to extract the most descriptive

terms, as well as to calculate the similarity between 2 terms. However, TF-IDF is based on the bag-of-words (BoW) model, which is why it disregards aspects such as text position, semantics and co-occurrence.

2.2 Word Embeddings

Word Embeddings are based on the approach of Harris Distributional Hypothesis from 1951, which states, that words that occur in the same contexts tend to have similar meanings.

A Word Embedding provides a word vector for every word. It takes a word and gives it a vector representation by extracting features from that word within the context that word appears in and assigns it a place within a vector space. Two similar word will occupy places which are close to each other within that vector space whereas words that are different, will have locations much further away from each other. This allows computation of distance calculation, which is why, for example, we're able to tell which word is similar to "small" in the same sense as "biggest" is similar to "big". We simply have to compute the vector X , which is equal to $Vector_{biggest} - Vector_{big} + Vector_{small}$ and then we find the word which is closest to X in the vector space, measured by cosine distance. [Efficient Estimation of Word Representations in Vector Space]

2.2.1 Word2Vec

Word2Vec is a "2-Model Architecture for computing continuous vector representations of words from very large dataset" [Efficient Estimation of Word Representations in Vector Space] that creates an n -dimensional vector space in which each word is represented as a vector. Word2Vecs 2 learning models are the CBoW and Skip-Gram-Model (Figure 3.1).

CBoW uses the context word to predict the target word. The input is a one-hot encoded vector. The weights between the input layer and the output layer can be represented by a $V \cdot N$ matrix W where each row of W is the N -dimensional vector representation v_W of the input word [word2vec Parameter Learning Explained]. The hidden-layer h is computed by multiplying the one-hot encoding vector of the input word w_I with the weight matrix W .

$$h = W_{(k, \cdot)}^T := v_{w_I}^T$$

Next we have another weight matrix $W' = w'_{ij}$ which is an $N \cdot V$ matrix. With these weights we can finally compute a score u_j for each word in the vocabulary

$$u_j = v_{w_j}^T h$$

where v'_{w_j} is the j -th column of the matrix W' . Afterwards we use *softmax*, which is a log-linear classification model, to obtain the posterior distribution of words.

$$p(w_j|w_I) = y_j = \frac{\exp(u_j)}{\sum_{j'=1}^V \exp(u_{j'})}$$

In contrast to the CBoW model, Skip-Gram uses the target word to predict the context words. The input is still a one-hot encoding vector, the hidden layers definition stays the same as in the CBoW model, each output is still using the same hidden layer to output matrix as in the CBoW model $p(w_{c,j} = w_{O,c}|w_I) = y_{c,j} = p(w_j|w_I) = y_j$ and the function for $u_j = u_{c,j}$ stays the same. However in the output layer, we are now outputting C multinomial distributions.

2.3 SVM

2.4 Multinomial Naive Bayes Classifier

The Naive Bayes classifier is based on Bayes' theorem, which comes from the probability calculus and describes the calculation of conditional probability. Each object in this classification approach is assigned to the class for which the highest probability was computed or for which the lowest costs arise in this assignment. The Multinomial Naive Bayes classifiers assumes that the position of the word does not matter, as well as that the feature probabilities $P(x_i|c_j)$ are independent given a class c . "The probability of a class value c given a test document d is computed as

$$P(c|d) = \frac{P(c) \prod_{w \in d} P(w|c)^{n_{wd}}}{P(d)}$$

where n_{wd} is the number of times word w occurs in document d , $P(w|c)$ is the probability of observing word w given class c , $P(c)$ is the prior probability of class c , and $P(d)$ is a constant that makes the probabilities for the different classes sum to one". [<https://link.springer.com/content/pdf/10.1007%2F11871637.pdf>]

2.5 Decision Trees and Random Forest Classifier

The Random Forest classifier is a classification technique that creates multiple decision trees from randomly selected subsets of training data. Each tree in this process may make a decision, these votes are then aggregated to determine the final class. According to Breiman [<https://www.stat.berkeley.edu/~breiman/randomforest2001.pdf>] the following algorithm should be used for each tree in the forest:

Algorithm 1 Random Forest

- 1: Take n bootstrap-samples.
These sample will be used as the training set for growing the tree
 - 2: Out of M characteristics of the training data, $m < M$ features are randomly selected at each node in the tree and the best split on these m features is used for splitting the node.
The subsequent selection of a feature from this set may be done by minimizing entropy
 - 3: The tree is fully expanded. No pruning is used
-

2.6 Logistic Regression Classifier

Logistic Regression is like Naive Bayes, a probabilistic classifier, and thus classifies by estimating the probability that the object belongs to a particular class. It can be derived analogously to the linear regression hypothesis,

$$h_{\Theta}(x) = \Theta^T x$$

but the logistic regression hypothesis generalizes the linear one to use the logistic function

$$h_{\Theta}(x) = g(\Theta^T x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

which results in

$$h_{\Theta}(x) = \frac{1}{1 + e^{-\Theta x}}$$

where $g(z)$ is the logistic function, also know as the *sigmoid function*.

Chapter 3

Data

3.1 Data Description

The given data, on which we want to build our model on, was provided by zenodo.org as part of SemEvals Task 4 [Link zu Task hinzufügen] and consists of 2 independent datasets, which in turn have been divided into a GroundTruth-, Training- and Validation set.

The first dataset, recognizable by the term 'byPublisher', reflects the publisher's general bias set forth by BuzzFeed journalists or MediaBiasFastCheck.com beforehand. It consists of a total of 750,000 items, of which 600,000 belong to the Training- and 150,000 to the Validation set.

In return, the second dataset, recognizable by the term 'byArticle', was scrapped by crowdsourcing at hand and therefore consists of only 645 items without a Validation set.

The GroundTruth dataset was provided as an XML File and consists of the features 'article url', 'labeled-by', 'id' and 'hyperpartisan'. In addition, the GroundTruth dataset scrapped 'byPublisher' contains the feature 'bias'.

- Article-url: Contains the article's url.
- Labeled-by: Reflects whether the respective article was labeled 'byPublisher' or 'byArticle'.
- Id: Allocates each article a unique id.
- Hyperpartisan: Displays whether the particular article has been labeled as hyperpartisan or not.
- Bias: Divides the publisher's bias into 'left', 'left-center', 'least', 'right-center' and 'right'.

The Training dataset was as well provided as an XML file and contains the contents of the website of the respective article. In addition, it consists of the features 'article title' 'published-at' and 'id'.

- Article title: Represents the articles title.
- Published-at: Specifies the published date.
- Id: A unique id, which is the same as the corresponding entry in the GroundTruth dataset.

The given Data has been cleaned in advance, therefore no additional steps were necessary.

The main focus of the datasets is on the Hyperpartisan feature, on which we want to classify the articles as this thesis progresses.

3.1.1 Dataset labelled by Publisher

As mentioned above, this dataset consists of a total of 750,000 articles and is divided into a training record consisting of 600,000 articles and a validation set consisting of 150,000 articles. Summarizing these two sets of data, a total of 375,000 were labelled as 'Hyperpartisan' and 375,000 were not – which corresponds to a 50:50 distribution. But even individually, this distribution does not change.




Figure 3.1: Distribution of as Hyperpartisan labelled articles by publisher


This dataset also includes the feature 'bias', which informs you about the general bias of the publisher. All 375,000 Hyperpartisan labelled all are assigned to either the left or right sectors, but none are right-centre, least or left-centre and are again 50:50 distributed.

The other 50% are split between the remaining bias, with 'Least' owning the largest share at 37%.


The publicity data is distributed over the years 1964-2018, with most of the data coming from 2012-2018.

3.1.2 Dataset labelled by Article

The dataset labelled by Article is a little different to the larger one labelled by publisher. Here the articles were individually labelled by hand. Accordingly, the




(a) Distribution of Bias




(b) Distribution of publishing years

distributions of this dataset are completely different. This becomes quiet obvious if we look at how the distribution of the Hyperpartisan labelled articles is here. Here



(c) Distribution of Hyperpartisan labelled articles



(d) Distribution of publishing years

we can see that there is no 50:50 distribution left. Only 36.9% were defined here as Hyperpartisan as shown in figure 2.4.

Moreover, in this dataset, the distribution of publication data is not mainly from the years 2012-2018, but in 2016-2018, with the largest number of articles dating back to 2017 at just under 60% as we can see in figure 2.5. Altogether all 645 articles date from the years 1902-2018.

3.2 Data Analysis

Hyperpartisan means "extremely partisan; extremely biased in favor of a political party." [definition]. This often materializes in relation to significant political events. As a result, in the following chapter, I will discuss in detail the direct link between the various features, especially the correlations between publication dates and label, as well as the connections between publisher and label.

Publisher	Amount
The Gateway Pundit	17
OpsLens	14
RealClearPolitics	13
New York Post	10
Salon	8

Table 3.1: Publishers who have published more than 7 Hyperpartisan Articles

Publisher	Bias	Amount
Fox Business	Left	96175
CounterPunch	Left	39832
Mother Jones	Left	36730
Truthdig	Left	25056
Daily Wire	Right	18570

Table 3.2: Publishers who have published more than 10.000 Hyperpartisan articles

The two tables 2.1 and 2.2 list the publisher who produced the highest amount of Hyperpartisan articles in the respective dataset. To keep track, only those publishers who have published more than 7 Hyperpartisan articles in the dataset 'labelled by article' and who have published more than 10.000 hyperpartisan articles in the dataset 'labelled by publisher' are included in the tables. However, a problem here is the dataset 'Labelled by Publisher', since this record, as previously described, has been labelled by the overall bias of the publisher. Meaning, for the further development, we can not include the publishers, since each article of a publishing house has the same label. What I would like to discuss later, however, is the connection between whether or not one of the publishers listed in Table 2.2 published more articles in important political years.

3.3 Data Preparation

In order to be able to work with the existing data in the further course of this project, several preprocessing steps are necessary. In the preprocessing phase of my bachelor's thesis, the data therefore went through the following steps:

1. Read the XML files.
2. Filter out the important information.
3. Merge the Groundtruth- and Training datasets into a csv file.
4. Remove special characters and stop words.
5. Tokenization and stemming of the datasets.

As a result, in the following section, I will go further into detail how I preprocessed my data.

3.3.1 Read the XML files

Since it is difficult to work with the given data in an XML format, it is necessary to bring them into a format with which it'll be easier to work with. The particular challenge hereby is the size of the dataset labelled by publisher. A standard algorithm for reading XML files is provided by python's DOM library ElementTree, called "ElementTree.parse". This method returns an ElementTree type, which "is a flexible container object, designed to store hierarchical data structures in memory" [<http://effbot.org/zone/element.htm>]. Meaning, that this library forms the entire model in the memory which can pose a problem with very large files, such as ours. As a substitute, I therefore use the method "ElementTree.iterparse", which can process XML documents in a streaming fashion, retaining in memory only the most recently examined parts of the tree.

3.3.2 Filter out the important information

As mentioned in Chapter 2.1 Data Description, the XML files include various features, which will play an important role in the further course of this work. It is now necessary to filter these features out of the XML format so they can be better worked with later. In order to be able to do this, we pass an algorithm that, by the iterparse method read content, which then passes through this algorithm in a double for-loop and checks for each item of an element in the content which "key" is currently processed. I then save this in an array, so that the features which have already been parsed, can be used later.

```
def parse_features(content, publisher):
    for event, elem in content:
        for key, value, in elem.items():
            if publisher:
                if key == 'id':
                    id_array_publisher.append(str(value))
                elif key == 'published-at':
                    published_at_array_publisher.append(value)
                elif key == 'title':
                    title_array_publisher.append(value)
            else:
                if key == 'id':
                    id_array_article.append(str(value))
                elif key == 'published-at':
                    published_at_array_article.append(value)
                elif key == 'title':
                    title_array_article.append(value)
        elem.clear()
```

3.3.3 Merge the Groundtruth- and Training datasets into a csv file

Since both, the Groundtruth- and Trainingdata contain important information, it is necessary to merge them into one file. I decided to use Python's library pandas to write the two files into one csv file. This, and especially Pandas, allows us to read the file more quickly as well as to access individual rows and columns of the

merged file in a targeted manner.

```
feature_extraction.parse_features(xml_training, publisher)
groundtruth_parser.parse_groundtruth(xml_gt, publisher)
content = content_parser.parse_content(content_training)

a_id = feature_extraction.get_id_array(publisher)
published = feature_extraction.get_published_at_array(publisher)
title = feature_extraction.get_title_array(publisher)
bias = groundtruth_parser.get_bias_array(publisher)
hyperpartisan = groundtruth_parser.get_hyperpartisan_array(publisher)

columns = {"ArticleID": a_id,
           "PublishedAt": published,
           "Title": title,
           "Bias": bias,
           "Content": content,
           "Hyperpartisan": hyperpartisan}

tf = Pd.DataFrame(columns)
tf = tf[['ArticleID', 'Published', 'Title', 'Bias', 'Content', 'Hyperpartisan']]
tf.to_csv('titlecsv', encoding='utf-8', index=False)
```

3.3.4 Remove special characters and stop words

Now that we have merged both files and written them into a csv file, we need to remove special characters and stop words. Especially the removal of stopwords is necessary since not all words presented in a document, such as auxiliary verbs, conjunctions and articles [TextClassificationUsingMachineLearning] are useful for training a classifier. It also decreases our corpus which makes it easier to classify later on. (Warum?)

```
stop = stopwords.words('english')

df['Content'] = df['Content'].apply(lambda x: ''.join([item for item in x.
split() if item not in stop]))
df['Title'] = df['Title'].apply(lambda x: ''.join([item for item in x.
split() if item not in stop]))
df['Content'] = df['Content'].map(lambda x: re.sub(r"^[a-zA-Z0-9]+", '', x))
df['Title'] = df['Title'].map(lambda x: re.sub(r"^[a-zA-Z0-9]+", '', x))
```

Here it becomes obvious that using pandas was a good choice, as we can now specifically access the 'Content' and 'Title' columns in order to perform this step of preprocessing on only these two and not the whole file.

3.3.5 Tokenization and Stemming of the datasets

After cleaning the dataset, the words are tokenized in order to convert them into numerical vectors so that a classifier is able to work with them. Tokenization is defined as "The process of demarcating and possibly classifying sections of a string of input characters. The resulting tokens are then passed on to some other form of processing. The process can be considered a sub-task of parsing input".

Stemming is the procedure of reducing the word to its grammatical (morpho-syntactic) root. The result is not necessarily a valid word of the language. For example, "recognized" would be converted to "recogniz". Still, the basic word almost always contains the very meaning of the word. Stemming is advantageous in that the algorithm used later now only has to fall back on a few different words and not many, all of which have the same meaning.

In order to implement Stemming and Tokenization, only 2 lines of python code are necessary, due to the Pandas dataframe.

```
df["Content"] = df["Content"].apply(nltk.word_tokenize)
df['Content'] = df['Content'].apply(lambda x: [stemmer.stem(y) for y in x
])
```

This will transform our output file in the following way:

Chapter 4

Methodology

The main process for Text Classification includes 7 steps of which the first 4 have already been performed in Chapter 2.

1. Read the Document.
2. Tokenize Text.
3. Stemming.
4. Delete Stopwords.
5. Vector Representation of the Text.
6. Feature Selection and/or Feature Transformation.
7. Learning Algorithm.

As already mentioned in Chapter "Introduction", we build our classifier model by using BERT-Embeddings. To compare how our model performs, in this chapter, I will discuss the classic methods and algorithms that I have used in order to achieve this comparison.

4.1 Important Algorithms

In order to avoid the repetition of algorithms used in the same context, in this section, I will explain some recurring algorithms.

Algorithm 2 Saving a trained model

```
import pickle
filename = 'finalized_model.sav'
pickle.dump(model, open(filename, 'wb'))
```

Algorithm 3 Loading a trained model

```
loaded_model = pickle.load(open(filename, 'rb'))
```

4.2 Vector Representation of the Text

In order for our classifier to be able to work with the text, we first need to transform our words into a feature vector representation. A document is a sequence of words [Text Categorization with Support Vector Machines] so a document can be presented by a One-Hot encoded vector, assigning the value 1 if the document contains the feature-word or 0 if the word does not appear in the document. [TextClassificationUsingmachineLearning]. However, using this technique for word representation, resolves in a $V \cdot V$ Matrix, as we have V-dimensional vector for each out of V words which can lead to huge memory issues. In addition this does not notion similarity between words. Therefore I will go into further detail for better approaches in the next 2 subchapters.

4.2.1 Term Frequency-Inverse Document Frequency

A comparative approach I used in the course of my Bachelor Thesis is Term Frequency-Inverse Document Frequency. By using TF-IDF, we're able to represent our word as a vector by assigning it weight which is computed thorough Term-Frequency divided by Inverse-Term-Frequency. Python's library *sklearn* provides two ways to implement TF-IDF without having to program TF and IDF by itself. In order to get a generally better overview, I will now explain the 2-step implementation, but also briefly explain how both steps can be combined in one.

Term-Frequency, as mentioned in chapter 1 – Principles, is a measure that denotes how often a term appears in a document. Inverse Document Frequency, on the other hand, reflects the importance of a term throughout the document corpus. To implemente the TF-IDF measure, sklearn provides the classes *CountVectorizer* and *TfidfTransformer* of the submodule *sklearn.feature_extraction.text*. Therefore, in the first step of our TF-IDF-implementation, both classes must be imported. As an input example, I use the first article of the "byArticle" labelled dataset.

```
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
```

```
df = pd.read_csv('/home/lstrauch/Bachelorarbeit/env/Data/
Preprocessed_ByArticle.csv', encoding='utf-8')
content= [df.Content[1]]
```

```
[ 'donald, trump, ran, many, braggadocios, largely, unrealistic, campaign,
  promises, one, promises, best, hugest, competent, infrastructure,
  president, united, states, ever, seen, trump, going, fix, every,
  infrastructure, problem, country, make, america, great, process,
  unless, brown, american, case, even, massive, natural, disaster, like,
  hurricane, maria, puerto, rico, debt, puerto, rican, citizens, ... ]'
```

In order to calculate our TF measure, we use the method *fit_transform()* of the class *CountVectorizer*, which we pass our document corpus as a parameter.

```
count_vect = CountVectorizer()
content_counts = count_vect.fit_transform(content)
```

fit_transform() learns a vocabulary dictionary of all tokens [sklearn Documentation], counts how many times a term *t* occurs in a document *d* and converts the text document into a token matrix.

```
<1x84 sparse matrix of type '<class_'numpy.int64'>'
with 84 stored elements in Compressed Sparse Row format>
```

Any term found during this procedure is assigned a unique index corresponding to a column of the regressive matrix.

```
content_counts.toarray()
```

```
array([[1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 2, 1, 1, 1, 1, 2,
       2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 3, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 3, 1, 1, 1, 1, 1, 1,
       2, 1, 1, 1, 1, 1, 1, 1, 3, 1, 1, 1, 1, 1, 2, 1, 1, 1]])
dtype=int64)
```

The calculation of the IDF-measure in the second step is similar to the calculation of the TF-measure. Again, the method *fit_transform()* is called. In contrast to the method of the *CountVectorizer*, we no longer pass our text document as a parameter, but the token matrix of *CountVectorizer*.

```
tfidf_transformer = TfidfTransformer()
content_tfidf = tfidf_transformer.fit_transform(content_counts)
```

```
array([[
0.0860663 , 0.0860663 , 0.17213259, 0.0860663 , 0.0860663 ,
0.0860663 , 0.0860663 , 0.0860663 , 0.0860663 , 0.0860663 ,
0.0860663 , 0.17213259, 0.0860663 , 0.0860663 , 0.0860663 ,
0.0860663 , 0.17213259, 0.0860663 , 0.0860663 , 0.0860663 ,
0.0860663 , 0.17213259, 0.17213259, 0.0860663 , ...]])
```

It should be noted that the shape of the finalized tfidf matrix depends on the document corpus. The first time the conversion is carried out, the two objects of the respective classes learn the respective vocabulary through the keyword "*fit*". If we want to convert a second text document to TF-IDF vectors afterwards, which is supposed to be used in connection with the already converted text document, we have to make sure to use the same *CountVectorizer* and the same *TfidfTransformer*. This is necessary because otherwise the error message "Dimension Mismatch" would occur in later prediction calls. Let's take the following example:

If we convert the two datasets labelled by article and publisher separately by the

method described above, we obtain the following dimensions:

- By Publisher: (600000, 708863)
- By Article: (645, 708863)

The first number in the vectors refers to the number of rows in the document and the second, how many columns the matrix contains. As we can clearly see, the two dimensions match because the documents were converted with the same fitted models.

```
def tf_idf_forTrainingAndTesting(x_train, x_test):
    count_vect = CountVectorizer()
    tfidf_transformer = TfidfTransformer()

    x_train_counts = count_vect.fit_transform(x_train)
    x_train_tfidf = tfidf_transformer.fit_transform(
        x_train_counts)

    x_test_counts = count_vect.transform(x_test)
    x_test_tfidf = tfidf_transformer.transform(x_test_counts)
```

If we would call the methods *fit_transform()* for both datasets, *CountVectorizer* and *TfidfTransformer* would be initialized each time, which would result in the following dimensions:

- By Publisher: (600000, 708863)
- By Article: (645, 11485)

Since it is not always possible or useful to convert the test data in the same step as the training data, the module *pickle* provides an opportunity to save the already initialized *CountVectorizer* and *TfidfTransformer* for later use on a new body of text.

```
pickle.dump(count_vect, open("/mypath/count_vect.pickle", "wb"))
pickle.dump(tfidf_transformer, open("/mypath/tfidf.pickle", "wb"))
```

This would make it possible to save the already fitted models.

```
pickle.dump(count_vect, open("/mypath/count_vect.pickle", "wb"))
pickle.dump(tfidf_transformer, open("/mypath/tfidf.pickle", "wb"))
```

Algorithm 4 pickle.dump

```
pickle.dump(count_vect, open("/mypath/count_vect.pickle", "wb"))
pickle.dump(tfidf_transformer, open("/mypath/tfidf.pickle", "wb"))
```

Term	Average TF-IDF Weight
trump	0.091165
bannon	0.065129
president	0.061829
kimmel	0.052415
money	0.051647
americans	0.049132
people	0.047342
obamacare	0.042512
wall	0.042117
control	0.039068

Table 4.1: Top 10 terms by average TF-IDF weight (By Article)

Terms	Average TF-IDF Weight
balls	0.057251
said	0.051183
medicare	0.047450
martin	0.046519
school	0.044264
university	0.043572
says	0.042079
trump	0.041235
carrier	0.040903
degree	0.040704

Table 4.2: Top 10 terms by average TF-IDF weight (By Publisher)

4.2.2 Word Embeddings

Word2Vec

For a further comparison and an approximation to the the later on used classification model, I did not only use TF-IDF as a Vector representation method as part of my Bachelor Thesis, but also Word Embeddings - especially the *Word2Vec* model. Word2Vec represents words as vectors. Unlike the TF-IDF method, however, not only word frequencies and -priorities are considered, but also the connection of individual words to others. Again, several methods of implementation exist. As part of my Bachelor Thesis, I decided to use the library *gensim* to implement my Word2Vec model. With *gensim* it is possible to do unsupervised semantic modelling from plain text. This makes it possible to implement a Word2Vec model using only a few lines of code without having to program Skip-Gram or CBoW yourself.

```
model = gensim.models.Word2Vec(vocab, min_count=10, window=10, size
                               =300, iter=10)
```

The algorithm above shows that the implementation of the model is straightforward, as it is pretty much the only step we need to program. By default, *gensim* uses CBoW which can be changed by adding the following parameter to the parameters list: *sg=1*. As for the other parameters, 18 more exist which can be viewed at "<https://radimrehurek.com/gensim/models/word2vec.html>", but I decided to focus only on the important ones. *Vocab* is our text corpus, which needs to be transformed into a list of tokenized sentences. *Size* determines the dimension of the word vectors, *window* the maximum distance between the current and predicted word within a sentence, *min_count* how often a word must occur to be included in the vocabulary and *iter* how many iterations should be performed on the corpus. What exactly happens here is that we train a neural network with a single hidden

By Article	Similarity	By Publisher	Similarity
president	0.9992413520812988	obama	0.5890584588050842
donald	0.9992144107818604	clinton	0.5385712385177612
election	0.9980630874633789	bannon	0.5277745127677917
presidential	0.9978925585746765	priebus	0.5117671489715576
campaign	0.9978247880935669	hillary	0.48596829175949097

Table 4.3: Most similar word to 'trump'

layer on which the model is trained to predict the current word based on the context. The resulting vector consists of several features that describe the target word. After the model has been trained it is possible to find out information like the similarity between 2 words. Table 4.3 shows the most similar words to "trump", which was assigned the highest TF-IDF weight.

```
# similarity between identical words
model.wv.similarity(w1="trump", w2="trump")
1,0

# similarity between two unrelated words
model.wv.similarity(w1="trump", w2="explosion")
-0.11690721
```

The functionality here is that the cosine similarity is calculated between 2 words in the vector space. As the range of the cosine similarity can go from $[-1, 1]$, words that are completely the same are assigned the value 1 and words that are not similar at all are given the value -1 . Our resulting word vectors now have the dimension defined in the parameter *size*. In order to be able to form features from them, I averaged the Word Embeddings of all words in a sentence.

```
def sent_vectorizer(sent, model):
    sent_vec = []
    numw = 0
    for w in sent:
        if numw == 0:
            sent_vec = model[w]
        else:
            sent_vec = np.add(sent_vec, model[w])
```

In order not to have to re-train the model every time, since this can take some time on large data sets, it is as well possible to save the trained model in order to access it again later.

```
word_vectors = model.wv
fname = "article_vectors.kv"
word_vectors.save(fname)
```

4.3 Feature Selection and/or Feature Transformation

4.4 Learning Algorithms

Classification is about predicting a particular outcome based on given training data. For this prediction, a classification algorithm processes the training data set, which consists of a set of features and the respective predictions. The algorithm attempts to discover relationships between given features of the instances and the associated classes to learn a function which makes it possible to predict the correct class based on the features of an instance. Thereafter, the algorithm receive a test dataset which it has not seen before. This dataset contains the same features as the training set but not the corresponding class names. With the previously learned function, the algorithm now assigns a class name to each instance of the test record.

Classic classification algorithms include *Multinomial Naive Bayes*, *Support Vector Machines*, *Random Forest* and *Logistic Regression*, which is why, in the following chapter, I will explain the basic procedure for implementing these algorithms. In addition, I will discuss the aspect of *Grid Search*, which gives us the optimal parameter assignment to a given set of data for these algorithms.

4.4.1 SVM

4.4.2 Multinomial Naive Bayes Classifier

The *MultinomialNB* class of the library *scikit-learn* implements the Naive Bayes algorithm for multinomial distributed data and is one of the two classic Naive Bayes variants used in text classification. Here, the distribution is parameterized by vectors $V_y = (Y_{y1}, \dots, Y_{yn})$ for each class y , where n is the size of the vocabulary and V_{yi} is the probability $P(x_i|y)$ of feature i appearing in a sample belonging to class y . The parameter V_y is estimated by a smoothed version of the maximum likelihood

$$\hat{V}_{yi} = \frac{N_{yi} + \alpha}{N_y + \alpha n}$$

where $N_{yi} = \sum_{x \in T} x_i$ is the number of times a feature i appears in a sample of class y in the training set T and $N_y = \sum_{i=1}^n N_{yi}$ is the total count of all features for class y . The smoothing priors $\alpha \geq 0$ accounts for features not present in the learning samples and prevents zero probabilities in further computations.

For implementing Multinomial NB, the first step is to define the classifier `clf = MultinomialNB()` which includes the parameters `alpha`, `fit_prior` and `class_prior`. `alpha` specifies α 's value with which smoothing should be made. `fit_prior` specifies whether the class probabilities should be learned in advance and `class_prior` specifies the prior probabilities of the classes. After the classifier has been defined, we train it by calling the method `fit()`, to which the training data X and the associated labels y are passed as parameters. This method then adjusts the Naive Bayes clas-

sifier according to X and y .

```
from sklearn.naive_bayes import MultinomialNB
clf = MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)
clf.fit(X, y)
```

4.4.3 Random Forest Classifier

4.4.4 Logistic Regression Classifier

Ehrenwörtliche Erklärung

Ich versichere, dass ich die beiliegende Master-/Bachelorarbeit ohne Hilfe Dritter und ohne Benutzung anderer als der angegebenen Quellen und Hilfsmittel angefertigt und die den benutzten Quellen wörtlich oder inhaltlich entnommenen Stellen als solche kenntlich gemacht habe. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen. Ich bin mir bewusst, dass eine falsche Erklärung rechtliche Folgen haben wird.

Mannheim, den 31.08.2014

Unterschrift