

# Hyperpartisan News Detection

Bachelor Thesis

presented by  
Larissa Strauch  
Matriculation Number 1518629

submitted to the  
Data and Web Science Group  
Prof. Dr. Ponzetto  
University of Mannheim

Juli 2019

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Problem Statement . . . . .	1
1.2	Contribution . . . . .	1
1.3	Related Work . . . . .	1
<b>2</b>	<b>Fundamentals</b>	<b>2</b>
2.1	Term Frequency-Inverse Document Frequency . . . . .	2
2.2	Word Embeddings . . . . .	3
2.2.1	Word2Vec . . . . .	3
2.3	Multinomial Naive Bayes Classifier . . . . .	4
2.4	Logistic Regression Classifier . . . . .	4
2.5	Decision Trees and Random Forest Classifier . . . . .	6
2.5.1	Decision Trees . . . . .	6
2.5.2	Random Forest Classifier . . . . .	7
2.6	Cross Validation . . . . .	7
2.7	Evaluation . . . . .	8
<b>3</b>	<b>Data</b>	<b>9</b>
3.1	Data Description . . . . .	9
3.1.1	Dataset labelled by Publisher . . . . .	10
3.1.2	Dataset labelled by Article . . . . .	10
3.2	Data Preparation . . . . .	11
3.2.1	File Parsing . . . . .	11
3.2.2	Information Filtering . . . . .	12
3.2.3	Combine Data . . . . .	12
3.2.4	Special Characters and Stop Word Removal . . . . .	12
3.2.5	Tokenization and Stemming of the datasets . . . . .	13
<b>4</b>	<b>Methodology</b>	<b>14</b>
4.1	Vector Representation of the Text . . . . .	14
4.1.1	Term Frequency-Inverse Document Frequency . . . . .	14
4.1.2	Word Embeddings . . . . .	16
4.2	Classification Algorithms . . . . .	17

<i>CONTENTS</i>	ii
4.2.1 Multinomial Naive Bayes Classifier . . . . .	18
4.2.2 Random Forest Classifier . . . . .	18
4.2.3 Logistic Regression Classifier . . . . .	19
4.2.4 Grid Search . . . . .	19
4.2.5 Training and Classification . . . . .	21
<b>5 Bidirectional Encoder Representations from Transformers</b>	<b>23</b>
<b>6 Evaluation</b>	<b>24</b>
<b>7 Conclusion</b>	<b>25</b>

# List of Algorithms

1	Random Forest . . . . .	7
2	Parse Ground-Truth File . . . . .	12
3	Merge Groundtruth- and Trainingdatasets . . . . .	12
4	Remove special characters and stop words . . . . .	13
5	Special Character and Stopword Removal with Pandas and NLTK . . . . .	13
6	TF-IDF . . . . .	15
7	Word2Vec with gensim . . . . .	16
8	sent_vectorizer . . . . .	17
9	Grid Search for Multinomial NB classifier . . . . .	20
10	Classifier Training . . . . .	21
11	Make Predictions . . . . .	21
12	Saving a trained model . . . . .	21
13	Saving a trained Word2Vec model . . . . .	22

# List of Figures

2.1	sigmoid function [2]	4
3.1	Hyperpartisan Distribution by Publisher	10
3.2	Hyperpartisan Distribution by Article	11
3.3	Publishing Years Distribution by Article	11

# List of Tables

2.1	Confusion Matrix . . . . .	8
3.1	Publishers with the highest proportion of Hypepartisan articles . .	10
4.1	Top 10 terms by TF-IDF weight . . . . .	16
4.2	Most similar word to 'trump' . . . . .	17
4.3	Grid Search Result of the Multinomial NB Classifier . . . . .	19

# Glossary

**CBoW** Continous Bag of Words. 3, 4, 16

**IDF** Inverse Term Frequency. 2, 13, 14

**TF** Term Frequency. 2, 13, 14

**TF-IDF** Term Frequency-Inverse Term Frequency. vi, 2, 13, 14, 16, 17

# **Chapter 1**

## **Introduction**

### **1.1 Problem Statement**

### **1.2 Contribution**

### **1.3 Related Work**



## Chapter 2

# Fundamentals

### 2.1 Term Frequency-Inverse Document Frequency

Term frequency (TF) is a measure that denotes how frequently a term  $t$  appears in the document  $d$ . One way to compute TF is [5]:

$$tf(t, d) = \frac{1 + \log(tf_{t,d})}{1 + \log(\max_{t \in d}(tf_{t,d}))}$$

where  $1 + \log(tf_{t,d})$  reflects how many times the term  $t$  appears in document  $d$  and  $1 + \log(\max_{t \in d}(tf_{t,d}))$  is the highest occurrence of any term in document  $d$ .

Inverse Document Frequency (IDF) points to the assumption that the informativeness of the term  $t$  is inversely proportional to the number of documents in the collection in which the term appears [5].

$$idf(t_i) = \log \frac{N}{|\{d \in D : t_i \in d\}|}$$

Where  $N$  is the total amount of documents in a document set and  $|\{d \in D : t_i \in d\}|$  is the amount how many times the term  $t$  appears in the document set.

To compute the weight for the term  $t_i$  within the document  $d_j$  we simply multiply the *TF* and *IDF* components:

$$w_{ij} = tf(t_i, d_j) \cdot idf(t_i)$$

Therefore, TF-IDF indicates how significant a word is to a document in a collection or corpus. It is regularly used as a weighting factor in Information Retrieval and Text Mining. *The TF-IDF value increases proportionally to the number of times a word appears in the document, but is offset by the frequency of the word in the corpus, which helps to control the fact that some words are usually more common than others.*

TF-IDF is easy to compute. In addition, it is possible to extract the most descriptive

terms, as well as to calculate the similarity between 2 terms. However, TF-IDF is based on the bag-of-words (BoW) model, which is why it disregards aspects such as text position, semantics and co-occurrence.

## 2.2 Word Embeddings

Word Embeddings are based on the approach of Harris' Distributional Hypothesis [9] from 1951, which states, that words that occur in the same contexts tend to have similar meanings.

A Word Embedding provides a word vector for each word. This is done by extracting features from that word within the context in which it appears and assigning it a place within the vector space. Two similar words will occupy locations near one another within this vector space, while words that differ will have positions much further apart. This makes it possible to calculate the distance calculation by computing cosine distance.

### 2.2.1 Word2Vec

Word2Vec is a "2-Model Architecture for computing continuous vector representations of words from very large dataset"[27] that creates an  $n$ -dimensional vector space in which each word is represented as a vector. Word2Vecs 2 learning models are the CBoW and Skip-Gram-Model.

CBoW uses the context word to predict the target word. The input is a one-hot encoded vector. The weights between the input layer and the output layer can be represented by a  $V \cdot N$  matrix  $W$  where "each row of  $W$  is the  $N$ -dimension vector representation  $w_v$  of the associated word of the input layer" [19]. The hidden-layer  $h$  is computed by multiplying the one-hot encoding vector of the input word  $w_I$  with the weight matrix  $W$  [19].

$$h = W^T x = W_{(k, \cdot)}^T := v_{w_I}^T$$

Next we have another weight matrix  $W' = w'_{ij}$  which is an  $N \cdot V$  matrix. With these weights we can finally compute a score  $u_j$  for each word in the vocabulary [19]

$$u_j = v_{w_j}'^T h$$

where  $v_{w_j}'$  is the  $j$ -th column of the matrix  $W'$ . Afterwards "we can use *softmax*, which is a log-linear classification model, to obtain the posterior distribution of words" [19].

$$p(w_j | w_I) = y_j = \frac{\exp(u_j)}{\sum_{j'=1}^V \exp(u_{j'})}$$

In contrast to the CBoW model, Skip-Gram uses the target word to predict the

context words. The input is still a one-hot encoding vector, the hidden layers definition stays the same as in the CBoW model, each output is still using the same hidden layer to output matrix as in the CBoW model  $p(w_{c,j} = w_{O,c}|w_I) = y_{c,j} = p(w_j|w_I) = y_j$  and the function for  $u_j = u_{c,j}$  stays the same [19]. However in the output layer, we are now outputting  $C$  multinomial distributions.

## 2.3 Multinomial Naive Bayes Classifier

The Naive Bayes classifier is based on Bayes' theorem [1], which comes from the probability calculus and describes the calculation of conditional probability. Each object in this classification approach is assigned to the class for which the highest probability was computed or for which the lowest costs arise in this assignment. The Multinomial Naive Bayes classifiers assumes that the position of the word does not matter, as well as that the feature probabilities  $Pr(t_i|c)$  are independent given a class  $c$ . The probability of a class value  $c$  given a test document  $t_i$  is computed as [12]

$$Pr(c|t_i) = \left(\sum_n f_{ni}\right)! \prod_n \frac{Pr(w_n|c)^{f_{ni}}}{f_{ni}!}$$

where  $f_{ni}$  is the number of times a word  $n$  occurs in document  $t$  and  $Pr(t_i|c)$  is the probability of observing word  $n$  given class  $c$ .

## 2.4 Logistic Regression Classifier




Figure 2.1: sigmoid function [2]

Logistic Regression is like Naive Bayes, a probabilistic classifier, and thus classifies by estimating the probability  $P(Y|X)$  that the object belongs to a particular class. It can be derived analogously to the linear regression model [8],

$$P(Y|X) = Pr(Y = 1) = X\beta$$

where  $X$  is the vector of predictors  $\{X_1, \dots, X_n\}$  and  $\beta$  is the model's parameter vector. In contrast to linear regression, in logistic regression only values between 0

and 1 are obtained, which can be attributed to the addition of the sigmoid function [3] (Figure 2.1).

$$p(\bar{x}_i) = \sigma(\bar{x}_i) = \frac{1}{1 + e^{-z_i}}$$

, where  $z_i = \log\left(\frac{p(\bar{x}_i)}{1-p(\bar{x}_i)}\right)$ .

So if the probability distribution is modelled with a sigmoid, the following is obtained [3]:

$$z_i = \log\left(\frac{p(\bar{x}_i)}{1-p(\bar{x}_i)}\right) = \log\left(\frac{\frac{1}{1+e^{-z_i}}}{1-\frac{1}{1+e^{-z_i}}}\right) = \log\left(\frac{1}{e^{-z_i}}\right) = z_i$$

This opens up the possibility of defining the probability that a sample belongs to a class [3]:

$$p(y|\bar{x}_i) = \sigma(\bar{x}_i, \bar{\Theta})$$

where  $\Theta$  is a single parameter vector. To find the optimal  $\Theta$  by which we get an  $p(y|\bar{x}_i)$  that is close to 0 or 1, we maximize the log-likelihood relying upon the output class, therefore the optimization problem can be expressed, utilizing the indicator notion as the minimization of the loss function [2]:

$$l(\Theta) = - \sum_i y_i \log(\sigma(z_i)) + (1 - y_i) \log(1 - \sigma(z_i))$$

This implies, if  $y = 0$ , the first term ends up 0 and the second  $\log(1 - \sigma(z_i))$ , which results in a log-likelihood of 0. In the event that  $y = 1$ , the second term ends up 0 and the first one corresponds to the log-likelihood of  $z$ . Along these lines, both cases are integrated into a solitary articulation.

A major issue in machine learning is the aspect of overfitting. Overfitting tends to adjust the model too much to the training data. This happens when a model learns the details in the training data so that the performance of the model is contrarily influenced by new data. This implies that noise or random fluctuations in the training data are recorded and learned as concepts by the model. The issue is, that these concepts do not apply to new data and negatively impact the generalization of the model. Therefore, there is the process of *regularization*, which is a form of regression that decreases the coefficient estimation to 0 and assumes that smaller weights generate simpler models and thus helps avoid overfitting.

Two of the most commonly used regularization techniques are *Lasso Regression (L1)* and *Ridge Regression (L2)*, which differ mainly in the penalty term.

The *L1* loss function minimizes the sum of the absolute differences between the target value and the estimated values. If input characteristics have weights closer to zero, this results in a sparse L1 standard. In the Sparse solution, most of the input features have zero weights and very few features have non-zero weights. The L1 regulation offers a function selection. This is done by associating insignificant input features with zero weights and useful features with a non-zero weight. *L1*

solves the regularized logistic regression by minimizing the following cost function:

$$\min_{w,c} \|w\|_1 + C \sum_{i=1}^n \log(\exp(1 + e^{-y_i x^T x_i}))$$

whereas L2 solves it the following way:

$$\min_{w,c} \frac{1}{2} w^T w + C \sum_{i=1}^n \log(1 + e^{-y_i w^T x_i})$$

, where  $C > 0$  is the penalty parameter,  $w$  is the vector,  $w^T$  is an additional dimension,  $x_i$  is an instance of the vector and  $y_i$  is an instance label pair  $\in \{-1, +1\}$  [7]. The L2 loss function is essentially about minimizing the sum of the square of the differences between the nominal value and the estimated values. L2 control forces the weights to small weights, but does not make them zero and leads to a not sparse solution. In addition, L2 is not robust against outliers, since square terms inflate the error differences of the outliers and the regularization term tries to correct them by punishing the weights.

## 2.5 Decision Trees and Random Forest Classifier

### 2.5.1 Decision Trees

Starting from the root node of a tree, a feature is evaluated from which a branch is subsequently selected. This process is repeated until the last node in the tree (leaf) is reached, which supplies the corresponding class. Different approaches have been developed in Decision Trees. One of the first is called *Iterative Dichotomizer (ID3)*. However, this algorithm had the disadvantage that explicit functions were necessary, which led to the development of C4.5. Like ID3, C4.5 manages continuous values. In return to ID3, C4.5 could furthermore manage summarized and decreed values and transform a tree into a set of conditional expressions. Nevertheless, the most recent development called CART (Classification and Regression Trees) was introduced, which allows the use of absolute and numeric values, as well as the non-use of standard sets. In addition, CART trees can be utilized for characterization and fallback assignments. [3]

The algorithm uses impurity measures to select a particular branch from a leaf. Among others, the two most common measures are *Gini Impurity* and *Cross Entropy Index* [2]. These impurity measures are applied to each candidate subset, and the resulting values are combined (e.g., averaged) to provide a measure of the quality of the split,

$$I_{Gini}(j) = \sum_i p(i|j)(1 - p(i|j))$$

$$I_{Cross-entropy}(j) = - \sum_i p(i|j) \log(p(i|j))$$

where  $j$  is a certain node,  $p(i|j)$  is the probability with  $i \in [1, n]$  associated with each class.

The *Gini* Index measures how often a randomly selected element from the set would be mislabelled if it were randomly chosen according to the distribution of labels in the subset. The minimum index (0,0) is achieved when all examples are classified into a solitary class.

”*Cross-Entropy* is based on information theory, and assumes null values only when samples belonging to a single class are present in a split, while it is maximum when there’s a uniform distribution among classes. This index is very similar to the Gini impurity, even though, more formally, the cross-entropy allows to select the split that minimizes the uncertainty about the classification, while the *Gini* impurity minimizes the probability of misclassification” [2].

### 2.5.2 Random Forest Classifier

The Random Forest classifier is a classification technique that creates multiple decision trees from randomly selected subsets of training data. Each tree in this process may make a decision, these votes are then aggregated to determine the final class. According to Breiman [4] the Random Forest algorithm is as follows:

---

#### Algorithm 1: Random Forest

---

- 1 Set the number of decision trees  $N_c$
  - 2 **for**  $i \leftarrow 1$  **to**  $N_c$  **do**
  - 3     Create a dataset  $D_i$  sampling with replacements from the original dataset  $X$
  - 4 Set the number of features to consider during each split  $N_f$
  - 5 Set an impurity measure
  - 6 Define an optimal maximum for each tree
  - 7 **for**  $i \leftarrow 1$  **to**  $N_c$  **do**
  - 8     Random Forest: Train the decision tree  $d_i(x)$  using the dataset  $D_i$  and selecting the best split among  $N_f$  features randomly sampled
  - 9     Extra-trees: Train the decision tree  $d_i(x)$  using the dataset  $D_i$  computing before each split  $n$  random thresholds and selecting the one that yield the least impurity
  - 10 Define an output function averaging the single outputs or employing a majority vote
- 

## 2.6 Cross Validation

*Cross Validation* is a model validation technique used to survey how the result of measurable statistical analysis generalizes into an independent dataset. The idea is to divide the entire dataset into a moving test and training set. The size of the test set is determined by the quantity of folds, so that at  $k$  emphases the test set covers the whole original dataset.

A round of Cross Validation comprises of separating a sample of data into corresponding subsets, performing the analysis on the training set, and validating the analysis on the test set. To decrease fluctuation, most strategies perform several rounds of *Cross Validation* utilizing various partitions and combine the validation result over the rounds to obtain an estimate of the predictive performance of the

model.

## 2.7 Evaluation

In order to evaluate how well a classifier works, several procedures exist. The ones used during the competition are *Accuracy*, *Recall*, *Precision* and *F1-Score*, for which the Confusion Matrix [25] (Table 2.1) forms the basis. Each column of the matrix represents the instances of a predicted class, while each row represents the instances of the actual class.

		Predicted Class	
		positive	negative
Actual Class	positive'	True Positives	False Negatives
	negative'	False Positives	True Negatives

Table 2.1: Confusion Matrix

*True Positive* means that the classifier predicted a class which actually corresponds to it, whereas *False Positive* means that a class was predicted that does not correspond to the actual class. In contrast, there are the *False Negatives*, where the classifier predicted a class as not belonging, although the instance actually belongs to it whereas *True Negative* means that the class was correctly classified as not belonging.

The four evaluation metrics are computed as follows [2]:

- Accuracy:  $\frac{TP+TN}{TP+TN+FP+FN}$   
Defines the correct classification to the total number of cases
- Precision:  $\frac{TP}{TP+FP}$   
Defines the correct classification of cases predicted to be positive
- Recall:  $\frac{TP}{TP+FN}$   
Defines the correct positive classification of cases that are actually positive
- F1-Score:  $2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$   
Defines the average of precision and recall, where an F1 value reaches its best value at 1 and worst at 0

## Chapter 3

# Data

### 3.1 Data Description

The given data, the model is build on, was provided by zenodo.org as part of SemEvals Task 4 [<https://pan.webis.de/semEval19/semEval19-web/>] and consists of 2 independent datasets, which in turn have been divided into a GroundTruth-, Training- and Validation set.

The first dataset, recognizable by the term 'byPublisher', reflects the publisher's general bias set forth by BuzzFeed journalists or MediaBiasFastCheck.com beforehand. It consists of a total of 750,000 items, of which 600,000 belong to the Training- and 150,000 to the Validation set.

In return, the second dataset, recognizable by the term 'byArticle', was scrapped by crowdsourcing at hand and therefore consists of only 645 items without a Validation set.

The GroundTruth-, Validation- and Training sets were each provided as XML documents. While the Training- and Validationset contain the articles, the GroundTruth file contains the attributes *article-url*, *labeled-by*, *id* and *hyperpartisan*. The main distinction between the by Article and by Publisher labelled datasets is that the by Publisher labelled GroundTruth dataset contains an additional attribute named *bias*. In this context, the *article-url* provides the URL of the article, the characteristic *labeled-by* reflects whether the article belongs to the publisher or article dataset, *id* represents a unique ID for the article, *hyperpartisan* reflects whether the article was labeled als Hyperpartisan or not and the additional attribute *bias* in the publisher record, states whether the article belongs to the "left", "left-center", "least", "right-center" or "right" sector.



### 3.1.1 Dataset labelled by Publisher




Figure 3.1: Hyperpartisan Distribution by Publisher

Publisher	Bias	Amount
Fox Business	Left	96175
CounterPunch	Left	39832
Mother Jones	Left	36730
Truthdig	Left	25056
Daily Wire	Right	18570

Table 3.1: Publishers with the highest proportion of Hyperpartisan articles

As mentioned above, this dataset consists of a total of 750,000 articles and is divided into a training record consisting of 600,000 articles and a validation set consisting of 150,000 articles. The main difference between this dataset and the by article labelled one is the type of classification. This is because these articles were not labelled as Hyperpartisan based on their content, but due to the publisher (Table 3.1). This means, that a publisher who has been designated as right or left by the MediaBiasFastCheck.com will naturally publish articles which are Hyperpartisan. This manifests itself in a different form of features, which becomes apparent in the further course of this thesis. This likewise influences the distribution of Hyperpartisan articles. Out of a total of 750,000 items, 375,000 assume the value 'True', while the remaining 375,000 have the value 'False' (Figure 3.1). Furthermore, the classification is expressed by the distribution of the additionally contained GroundTruth attribute *Bias*, which informs about the general bias of the publisher. All 375,000 Hyperpartisan labelled all are assigned to either the left or right sectors, but none are right-centre, least or left-centre and are again 50:50 distributed. The other 50% are split between the remaining bias, with 'Least' owning the largest share at 37%. The publicity data is distributed over the years 1964-2018, with most of the data coming from 2012-2018.

### 3.1.2 Dataset labelled by Article

As already described in the previous section, the distribution of the by Article dataset is different from the by Publisher one. The distribution is no longer 50:50, but only 36.9% were classified as Hyperpartisan (Figure 3.2). Moreover, in this dataset, the distribution of publication data is not mainly from the years 2012-2018, but in 2016-2018, with the largest number of articles dating back to 2017 at just under 60% as we can see in figure 3.3. Altogether all 645 articles date from the years 1902-2018.




Figure 3.2: Hyperpartisan Distribution by Article




Figure 3.3: Publishing Years Distribution by Article

## 3.2 Data Preparation

In order to be able to work with the existing data in the further course of this project, several preprocessing steps are necessary. In the preprocessing phase of my bachelor's thesis, the data therefore went through the following steps:

1. File Parsing.
2. Information Filtering.
3. Combine Data.
4. Special Characters and Stop Word Removal.
5. Tokenization and Stemming.

As a result, in the following section, I will go further into detail how I preprocessed my data.

### 3.2.1 File Parsing

Since it is difficult to work with the given data in an XML format, the first challenge is to convert these files into a format which allows work with them more easily. The particular challenge here is the size of the dataset labelled by publisher. A standard algorithm for reading XML files is provided by python's DOM library *ElementTree*, called *ElementTree.parse*[6]. This method returns an *ElementTree* type, which "is a flexible container object, designed to store hierarchical data structures in memory" [ <http://effbot.org/zone/element.htm>]. Meaning, that this library forms the entire model in the memory which can pose a problem with very large files, such as ours. As a substitute, I therefore use the method *ElementTree.iterparse*, which can process XML documents in a streaming fashion, retaining in memory only the most recently examined parts of the tree[13].

### 3.2.2 Information Filtering

As mentioned in Chapter 2.1 Data Description, the XML files include various features, which is why it is necessary to extract these features from the XML files. In order to do this, an algorithm is passed the, by the *iterparse* method read content, which then passes through this algorithm in a double for-loop and checks for each item of an element in the content which "key" is currently processed. I then save this in an array, so that the features which have already been parsed, can be used later.

---

**Algorithm 2:** Parse Ground-Truth File
 

---

```

Input: Ground-Truth File
Output: Trained Classifier
1 for event, elem in content do
2   for key, value in elem.items() do
3     if key == 'id' then
4       id_array.append(str(value))
5     else if key == 'published-at' then
6       published_at_array.append(value)
7     else if key == 'title' then
8       title_array.append(value)

```

---

### 3.2.3 Combine Data

Since both, the Groundtruth- and Trainingdata contain important information, it is necessary to merge them into one file. I decided to use Python's library *pandas* [14] to combine both files into a single one. This, and especially Pandas, allows us to read the file more quickly as well as to access individual rows and columns of the merged file in a targeted manner.

---

**Algorithm 3:** Merge Groundtruth- and Trainingdatasets
 

---

```

1 content = content_parser.parse_content(content_training)
2 a_id = feature_extraction.get_id_array()

3 published = feature_extraction.get_published_at_array()
4 title = feature_extraction.get_title_array()
5 bias = groundtruth_parser.get_bias_array()
6 hyperpartisan = groundtruth_parser.get_hyperpartisan_array()

7 columns = {"ArticleID": a_id, "PublishedAt": published, "Title": title, "Bias": bias, "Content":
            content "Hyperpartisan": hyperpartisan}

8 tf = Pd.DataFrame(columns)
9 tf = tf[['ArticleID', 'Published', 'Title', 'Bias', 'Content', 'Hyperpartisan']]
10 tf.to_csv(titlecsv, encoding='utf-8', index=False)

```

---

### 3.2.4 Special Characters and Stop Word Removal

Since the GroundTruth- and Trainingdatasets have been combined into a single file, the next step is to remove special characters and stop words. Especially the removal of stopwords is necessary since not all words presented in a document, such

as auxiliary verbs, conjunctions and articles [10] are useful for training a classifier.

---

**Algorithm 4:** Remove special characters and stop words
 

---

```

1 stop = stopwords.words('english')
2 df.Content = df.Content.apply(lambda x: ".join([item for item in x.split() if item not in stop])")
3 df.Content = df.Content.map(lambda x: re.sub(r"[^ a-zA-Z0-9]+", "", x))
```

---

### 3.2.5 Tokenization and Stemming of the datasets

After cleaning the dataset, the words are tokenized in order to convert them into numerical vectors so that a classifier is able to work with them. Tokenization is defined as "The process of demarcating and possibly classifying sections of a string of input characters". The resulting tokens are then passed on to some other form of processing. The process can be considered a sub-task of parsing input.

Stemming is the procedure of reducing the word to its grammatical root. The result is not necessarily a valid word of the language. For example, "recognized" would be converted to "recogniz". Still, the basic word almost always contains the very meaning of the word. Stemming is advantageous in that the algorithm used later now only has to fall back on a few different words instead of many, all of which have the same meaning.

In order to implement Stemming and Tokenization, *NLTK* [26] package provides two functions, due to which only 3 lines of code are necessary.

---

**Algorithm 5:** Special Character and Stopword Removal with Pandas and NLTK
 

---

```

1 import nltk
2 df.Content = df.Content.apply(nltk.word_tokenize)
3 df.Content = df.Content.apply(lambda x: [stemmer.stem(y) for y in x])
```

---

## Chapter 4

# Methodology

The primary procedure in Text Classification consist of 7 steps [10]. These include reading the dataset, tokenization, stemming, removing stop words, represent the text using vectors, feature extraction and selection, as well as applying classification algorithms. Since the first five steps have already been covered in Chapter 2, I will discuss the remaining three in this section. Especially, step 7 – applying classification algorithms – plays an important role. As already mentioned in Chapter 1 - Introduction, we build the classifier using BERTs-Embeddings. In order to compare how the model performs, I will therefore focus on classical text classification algorithms.

### 4.1 Vector Representation of the Text

In order for the classifier to be able to work with the text, the first step is to transform the words into a feature vector representation. A document is a sequence of words [11] so a document can be presented by a One-Hot encoded vector, assigning the value 1 if the document contains the feature-word or 0 if the word does not appear in the document [10]. However, using this technique for word representation, resolves in a  $V \cdot V$  Matrix, as we have  $V$ -dimensional vector for each out of  $V$  words which can lead to huge memory issues. In addition this does not notion similarity between words. Therefore I will go into further detail for better approaches in the next 2 subchapters.

#### 4.1.1 Term Frequency-Inverse Document Frequency

A comparative approach I used in the course of my Bachelor Thesis is Term Frequency - Inverse Document Frequency. Using TF-IDF allows to represent a word as a vector by assigning it weight which is computed through Term-Frequency multiplied with Inverse-Term-Frequency. Table 4.1 shows the 10 terms which have been assigned the highest TF-IDF weights. Python's library *scikit-learn* [16] provides two ways to implement TF-IDF without having to program TF and IDF by

**Algorithm 6:** TF-IDF**Input:** Text Corpus

---

```

1 count_vect = CountVectorizer()
2 content_counts = count_vect.fit_transform(corpus)

3 tfidf_transformer = TfidfTransformer()
4 content_tfidf = tfidf_transformer.fit_transform(content_counts)

```

---

itself. In order to get a generally better overview, I will now explain the 2-step implementation, but note that the class *sklearn.feature\_extraction.text.TfidfVectorizer* enables implementation in just one.

Term-Frequency, as mentioned in "Chapter 1 – Principles", is a measure that denotes how often a term appears in a document. Inverse Document Frequency, on the other hand, reflects the importance of a term throughout a document corpus. To implement the TF-IDF measure, *scikit-learn* provides the classes *CountVectorizer* and *TfidfTransformer* of the submodule *sklearn.feature\_extraction.text* [21]. As an input example, I use the first article of the "byArticle" labelled dataset.

In order to calculate the TF measure, the method *fit\_transform()* of the class *CountVectorizer* can be used, which the document corpus is passed as a parameter.

*fit\_transform()* learns a vocabulary dictionary of all tokens and then counts how many times a term  $t$  occurs in a document  $d$  and converts the text document into a token matrix. The calculation of the IDF-measure in the second step is similar to the calculation of the TF-measure. Again, the method *fit\_transform()* is called. In contrast to the method of the *CountVectorizer*, the text document is no longer passed as a parameter, but the token matrix of *CountVectorizer*.

It should be noted that the shape of the finalized TF-IDF matrix depends on the document corpus. The first time the conversion is carried out, the two objects of the respective classes learn the respective vocabulary through the keyword "fit\_". If a second text document should be converted to TF-IDF vectors afterwards, which is supposed to be used in connection with the already converted text document, it is necessary to use the same *CountVectorizer* and the same *TfidfTransformer*. This is, because otherwise the error message "Dimension Mismatch" would occur in later prediction calls. Let's take the following example:

If the two datasets labelled by article and publisher are converted with the same *CountVectorizer* and *TfidfTransformer*, we obtain the following dimensions:

- By Publisher: (600000, 708863)
- By Article: (645, 708863)

If the two methods *fit\_transform()* would be called for both datasets, *CountVectorizer* and *TfidfTransformer* would be initialized each time, which would result in the following dimensions:

By Article	TF-IDF Weight	By Publisher	TF-IDF Weight
trump	0.091165	balls	0.057251
bannon	0.065129	said	0.051183
president	0.061829	medicare	0.047450
kimmel	0.052415	martin	0.046519
money	0.051647	school	0.044264
americans	0.049132	university	0.043572
people	0.047342	says	0.042079
obamacare	0.042512	trump	0.041235
wall	0.042117	carrier	0.040903
control	0.039068	degree	0.040704

Table 4.1: Top 10 terms by TF-IDF weight

- By Publisher: (600000, 708863)
- By Article: (645, 11485)

#### 4.1.2 Word Embeddings

##### Word2Vec

For a further comparison and an approximation to the later on used classification model, I did not only use TF-IDF as a Vector representation method as part of my Bachelor Thesis, but also Word Embeddings - especially the *Word2Vec* model. Word2Vec represents words as vectors. Unlike the TF-IDF method, however, not only word frequencies and -priorities are considered, but also the connection of individual words to others. Again, several methods of implementation exist. As part of my Bachelor Thesis, I decided to use the library *gensim* [18] to implement my Word2Vec model. With *gensim* it is possible to do unsupervised semantic modelling from plain text. This makes it possible to implement a Word2Vec model using only a few lines of code without having to program Skip-Gram or CBoW yourself.

Algorithm 7 shows that the implementation of the model is straightforward, as it

---

##### Algorithm 7: Word2Vec with gensim

---

```
1 model = gensim.models.Word2Vec(vocab, min_count=10, window=10,
    size=300, iter=10)
```

---

is pretty much the only step we need to program. By default, *gensim* uses CBoW which can be changed by adding the following parameter to the parameters list: *sg=1*. As for the other parameters, 18 more exist which can be viewed at <https://radimrehurek.com/gensim/models/word2vec.html>, but I decided to focus only on the important ones. *Vocab* is our text corpus, which needs to be transformed into a list of tokenized sentences. *Size* determines the dimension of the word vectors, *window* the maximum distance between the current and predicted

By Article	Similarity	By Publisher	Similarity
president	0.9992413520812988	obama	0.5890584588050842
donald	0.9992144107818604	clinton	0.5385712385177612
election	0.9980630874633789	bannon	0.5277745127677917
presidential	0.9978925585746765	priebus	0.5117671489715576
campaign	0.9978247880935669	hillary	0.48596829175949097

Table 4.2: Most similar word to 'trump'

word within a sentence, *min\_count* how often a word must occur to be included in the vocabulary and *iter* how many iterations should be performed on the corpus. What exactly happens here is that a neural network with a single hidden layer is trained to predict the current word based on the context. The resulting vector consists of several features that describe the target word.

After the model has been trained it is possible to get information about the similarity of two words by calling the method *model.wv.similarity(word<sub>1</sub>, word<sub>2</sub>)*. Table 4.2 shows the 6 most similar words to "trump" of both datasets. This is possible by calculating the cosine similarity of two words in the vector space. As the range of the cosine similarity can go from [-1 to 1], words that are completely the same are assigned the value 1 and words that are not similar at all are given the value -1.

The resulting word vectors now have the dimension defined in the parameter *size*. In order to be able to form features from them, I averaged the Word Embeddings of all words in a sentence (algorithm 8).

---

**Algorithm 8:** sent\_vectorizer

---

**Input:** text corpus, trained model

```

1 sent_vec = []
2 numw = 0
3 for w in sent do
4     if numw == 0 then
5         sent_vec = model[w]
6     else
7         sent_vec = np.add(sent_vec, model[w])
8 return np.asarray(sent_vec) / numw

```

---

## 4.2 Classification Algorithms

Classification is about predicting a particular outcome based on given training data. For this prediction, a classification algorithm processes the training data set, which consists of a set of features and the respective predictions. The algorithm attempts



to discover relationships between given features of the instances and the associated classes to learn a function which makes it possible to predict the correct class based on the features of an instance. Thereafter, the algorithm receive a test dataset which it has not seen before. This dataset contains the same features as the training set but not the corresponding class names. With the previously learned function, the algorithm now assigns a class name to each instance of the test record.

Classic classification algorithms include *Multinomial Naive Bayes*, *Support Vector Machines*, *Random Forest* and *Logistic Regression*, which is why, in the following chapter, I will explain the basic procedure for implementing these algorithms. In addition, I will discuss the aspect of *Grid Search*, which gives us the optimal parameter assignment to a given set of data for these algorithms.

### 4.2.1 Multinomial Naive Bayes Classifier

The `sklearn.naive_bayes.MultinomialNB` class of the library *scikit-learn* implements the Naive Bayes algorithm for multinomial distributed data [24]. Here  $\Theta$  is estimated by a smoothed version of the maximum likelihood

$$\hat{\Theta}_{yi} = \frac{N_{yi} + \alpha}{N_y + \alpha n}$$

where  $N_{yi} = \sum_{x \in T} x_i$  is the number of times a feature  $i$  appears in a sample of class  $y$  in the training set,  $N_y = \sum_{i=1}^n N_{yi}$  is the total count of all features for class  $y$  and  $\alpha$  incorporates the smoothing parameter due to which zero probabilities will be prevented.

The *MultinomialNB* classifier includes the parameters *alpha*, *fit\_prior* and *class\_prior*, where *alpha* specifies the smoothing value, *fit\_prior* specifies whether the class probabilities should be learned in advance and *class\_prior* specifies the prior probabilities of the classes.

```
from sklearn.naive_bayes import MultinomialNB
MultinomialNB(alpha=0.5, fit_prior=True)
```

### 4.2.2 Random Forest Classifier

As described in Chapter 2.6.2, the Random Forest Classifier is a classification technique that creates multiple decision trees from randomly selected subsets of training data. For implementing the Random Forest Classifier, *scikit-learn* provides the class `sklearn.ensemble.RandomForestClassifier` [20]. Unlike the original publication [4], the *scikit-learn* classifier determines the final class by averaging the probabilistic forecasts as opposed to having each classifier vote in favour of a solitary class.

The Random Forest classifier includes 17 parameters, of which I have included 6 in my classification.

```
from sklearn.ensemble import RandomForestClassifier
RandomForestClassifier(min_samples_leaf=1, n_estimators=500, criterion='
    gini', max_depth=15, max_features='auto', bootstrap=True)
```

where *n\_estimators* determines the number of trees in the forest, *criterion* which impurity measure to utilize, *max\_depth* determines the maximum depth of a tree, *min\_samples\_leaf* determines the base number of samples to be at a leaf node, *max\_features* the quantity of features that must be considered in the search for the best partition and *bootstrap* indicates whether bootstrap models are utilized when making trees.

### 4.2.3 Logistic Regression Classifier

The Logistic Regression classifier is a model for classification, where the probabilities depicting the potential results of a solitary class are modelled utilizing a logistic function. The provided classifier of *scikit-learn* offers the possibility not only to classify binary, but also One-vs-Rest and multinomial. There are several solvers available for this, yet since our classification problem only refers to binary classification, I will only discuss those aspects of the logistic regression classifier in the following section.

Scikit learns Logistic Regression Classifier provides L1, L2 and Elastic-Net Regularization [22].

A solvers, *liblinear*, *newton-cg*, *lbfgs*, *sag* and *saga* are available. *Scikit-learn* points out, that *liblinear* is a good algorithm for small datasets, whereas *sag* and *saga* are faster for large datasets. Also, *newton-cg*, *lbfgs* and *sag* only handle L2 penalty, whereas *liblinear* also handles L1 penalty and *saga* in addition *elasticnet*. *Liblinear* uses a coordinate descent algorithm, the *sag* solver a Stochastic Average Gradient descent, *saga* is a variant of *sag* and therefor supports the non-smooth penalty L1 and *elasticnet* and the *lbfgs* solver is an optimization algorithm that approximates the Broyden–Fletcher–Goldfarb–Shanno algorithm

### 4.2.4 Grid Search

Parameter	Result by Article	Result by Publisher
alpha	0.7	0.5
fit_prior	False	True

Table 4.3: Grid Search Result of the Multinomial NB Classifier

A significant perspective in machine learning is the aspect of hyperparameters. These are parameters that are not determined by the learning algorithm, but those that must be determined beforehand. In the above mentioned algorithms, for example, these are the parameters to be passed to the classifier. For instance, which size *n\_estimators* in the *Random Forest* classifier should assume. To solve this issue the

algorithm *GridSearch* exists. This algorithm is used to find the optimal hyperparameters of a model, which then leads to the most accurate predictions.

The procedure is as follows: A set of parameters is defined with which GridSearch trains the given classifier for all possible combinations and measures the performance by cross-validation. This ensures that our trained model has received the most samples from our training data set.

A simple implementation of GridSearch is provided by *scikit-learn* through the class *sklearn.model\_selection.GridSearchCV* [23]. This class evaluates all possible combinations of parameters when calling the method *fit()* and keeps the best combination.

The parameters that can be passed to the class include *estimator*, which specifies the classifier, *param\_grid*, which is the parameter set, *scoring*, which specifies the measure to evaluate the test set, and *cv* how many cross-validation splits to perform. Algorithm 8 shows a practical example of how to use *GridSearchCV* while table 4.4 shows the corresponding results.

---

**Algorithm 9:** Grid Search for Multinomial NB classifier

---

```

1 from sklearn.model_selection import GridSearchCV
2 from sklearn.naive_bayes import MultinomialNB
3 parameter_candidates = {'alpha': np.linspace(0.5, 1.5, 6), 'fit_prior':
  [True, False]}
4 clf = GridSearchCV(estimator=MultinomialNB(),
  param_grid=parameter_candidates, scoring='accuracy', cv=10)
5 clf.fit(features, labels)
```

---

### 4.2.5 Training and Classification

Classifiers are algorithms which map input data to a specific type of category. Using the classifier consists of two basic steps: training and classification. Training is the process of recording content known to belong to certain class and creating a classifier based on that known content. Classification is the process by which a classifier is created with such a training content and is run on unknown content to determine the class affiliation for the undefined dataset. The above mentioned

---

**Algorithm 10:** Classifier Training
 

---

**Input:** Classifier model, Features, Labels

**Output:** Trained Classifier

```
1 classifier = model
2 classifier.fit(features, labels)
```

---

classifiers all follow the same scheme of classification (algorithm 10). To train these classifiers, the method *fit()* is called at the defined classification algorithm, to which the training data and its associated labels are passed on. Depending on the algorithm and the size of the data set, this can take different lengths of time. For example, the *RandomForest* Classifier classifies the by Article labelled data within X minutes, whereas it needs Y minutes for the by Publisher labelled data. Afterwards it is possible to make predictions with the trained classification. This is done by calling the *predict()* method to which the test data is passed (algorithm 11).

---

**Algorithm 11:** Make Predictions
 

---

**Input:** Trained model, testset

```
1 trained_model.predict(testset)
```

---

Finally, I consider it important to mention that it is possible to save trained classifiers and models in order to access them later on without having to train a classifier each time. For this purpose the module *pickle* of Python's standard library is available (algorithm 12), as well as for word vector models, the module *models.keyedvectors* (algorithm 13) of the library *gensim*. The *pickle* module implements binary protocols for serializing and deserializing a Python object structure, whereas *models.keyedvectors* implements word vectors and their similarity look-ups.

---

**Algorithm 12:** Saving a trained model
 

---

**Input:** Trained model, Final filename

```
1 filename = file pickle.dump(trained model, filename)
```

---

”The pickle module implements binary protocols for serializing and de-serializing

a Python object structure. 'Pickling' is the process whereby a Python object hierarchy is converted into a byte stream, and "unpickling" is the inverse operation, whereby a byte stream (from a binary file or bytes-like object) is converted back into an object hierarchy" [17].

---

**Algorithm 13:** Saving a trained Word2Vec model

---

**Input:** Word2Vec model, Final filename with ending "kv"

---

```
1 word_vectors = model.wv
2 filename = "vectors.kv"
3 word_vectors.save(filename)
```

---

*KeyedVectors* "is essentially a mapping between entities and vectors. Each entity is identified by its string id, so this is a mapping between a string and a 1D numpy array. The entity typically corresponds to a word, but for some models, the key can also correspond to a document, a graph node etc. To generalize over different use-cases, this module calls the keys entities. Each entity is always represented by its string id, no matter whether the entity is a word, a document or a graph node" [15].

## **Chapter 5**

# **Bidirectional Encoder Representations from Transformers**

## **Chapter 6**

# **Evaluation**

## **Chapter 7**

## **Conclusion**



# Bibliography

- [1] Mr. Bayes and Mr. Price. An essay towards solving a problem in the doctrine of chances. by the late rev. mr. bayes, f. r. s. communicated by mr. price, in a letter to john canton, a. m. f. r. s. *Philosophical Transactions (1683-1775)*, 53:370–418, 1763.
- [2] Giuseppe Bonaccorso. *Machine learning algorithms : reference guide for popular algorithms for data science and machine learning*. Birmingham, England ; Mumbai, India : Packt, 2017.
- [3] Giuseppe Bonaccorso. *Machine learning algorithms : popular algorithms for data science and machine learning*. Birmingham ; Mumbai : Packt Publishing, 2018.
- [4] Leo Breiman. Random forests. Technical report, University of California, January 2001.
- [5] Prabhakar Raghavan Christopher D. Manning and Hinrich Schütze. *Introduction to Information Retrieval*. Cambridge University Press., 2008.
- [6] Elementree. <https://github.com/python/cpython/blob/3.7/Lib/xml/etree/ElementTree.py>.
- [7] Rong-En Fan, Kai-Wei Chang, Cho-Jui Hsieh, Xiang-Rui Wang, and Chih-Jen Lin. Liblinear: A library for large linear classification. *J. Mach. Learn. Res.*, 9:1871–1874, June 2008.
- [8] Frank E. Harrell. *Binary Logistic Regression*, pages 219–274. Springer International Publishing, Cham, 2015.
- [9] Zellig S. Harris. Distributional structure. *WORD*, 10(2-3):146–162, 1954.
- [10] Emmanouil Ikonomakis, Sotiris Kotsiantis, and V Tampakas. Text classification using machine learning techniques. *WSEAS transactions on computers*, 4:966–974, 08 2005.
- [11] Thorsten Joachims. Text categorization with support vector machines: Learning with many relevant features. Technical report, Universität Dortmund, 1997.

- [12] Ashraf M. Kibriya, Eibe Frank, Bernhard Pfahringer, and Geoffrey Holmes. Multinomial naive bayes for text categorization revisited. In Geoffrey I. Webb and Xinghuo Yu, editors, *AI 2004: Advances in Artificial Intelligence*, pages 488–499, Berlin, Heidelberg, 2005. Springer Berlin Heidelberg.
- [13] Class `iterparse`. <https://lxml.de/api/lxml.etree.iterparse-class.html>.
- [14] Wes McKinney et al. Data structures for statistical computing in python. In *Proceedings of the 9th Python in Science Conference*, volume 445, pages 51–56. Austin, TX, 2010.
- [15] `models.keyedvectors` – store and query word vectors. <https://radimrehurek.com/gensim/models/keyedvectors.html>.
- [16] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [17] `pickle` — python object serialization. <https://docs.python.org/3/library/pickle.html>.
- [18] Radim Řehůřek and Petr Sojka. Software Framework for Topic Modelling with Large Corpora. In *Proceedings of the LREC 2010 Workshop on New Challenges for NLP Frameworks*, pages 45–50, Valletta, Malta, May 2010. ELRA. <http://is.muni.cz/publication/884893/en>.
- [19] Xin Rong. word2vec parameter learning explained. *CoRR*, abs/1411.2738, 2014.
- [20] `sklearn.ensemble.randomforestclassifier`. <https://github.com/scikit-learn/scikit-learn/blob/7813f7efb/sklearn/ensemble/forest.py#L758>.
- [21] `sklearn.feature_extraction.text`. [https://github.com/scikit-learn/scikit-learn/blob/7813f7efb5b2012412888b69e73d76f2df2b50b6/sklearn/feature\\_extraction/text.py](https://github.com/scikit-learn/scikit-learn/blob/7813f7efb5b2012412888b69e73d76f2df2b50b6/sklearn/feature_extraction/text.py).
- [22] `sklearn.linear_model`. [https://github.com/scikit-learn/scikit-learn/blob/7813f7efb/sklearn/linear\\_model/logistic.py#L1202](https://github.com/scikit-learn/scikit-learn/blob/7813f7efb/sklearn/linear_model/logistic.py#L1202).
- [23] `sklearn.model_selection.gridsearchcv`. [https://github.com/scikit-learn/scikit-learn/blob/7813f7efb/sklearn/model\\_selection/\\_search.py#L828](https://github.com/scikit-learn/scikit-learn/blob/7813f7efb/sklearn/model_selection/_search.py#L828).

- [24] `sklearn.naive_bayes.multinomialnb`. [https://github.com/scikit-learn/scikit-learn/blob/7813f7efb/sklearn/naive\\_bayes.py#L636](https://github.com/scikit-learn/scikit-learn/blob/7813f7efb/sklearn/naive_bayes.py#L636).
- [25] Stephen V. Stehman. Selecting and interpreting measures of thematic classification accuracy. *Remote Sensing of Environment*, 62(1):77 – 89, 1997.
- [26] Ewan Klein Steven Bird and Edward Loper. *Natural Language Processing with Python – Analyzing Text with the Natural Language Toolkit*. O'Reilly Media, 2009.
- [27] Greg Corrado Tomas Mikolov, Kai Chen and Jeffrey Dean. Efficient estimation of word representations in vector space. *arXiv.org*, January 2013.

## **Ehrenwörtliche Erklärung**

Ich versichere, dass ich die beiliegende Master-/Bachelorarbeit ohne Hilfe Dritter und ohne Benutzung anderer als der angegebenen Quellen und Hilfsmittel angefertigt und die den benutzten Quellen wörtlich oder inhaltlich entnommenen Stellen als solche kenntlich gemacht habe. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen. Ich bin mir bewusst, dass eine falsche Erklärung rechtliche Folgen haben wird.

Mannheim, den 31.08.2014

Unterschrift